

Lecture 3: C++ Thread Programming

Tsung-Wei (TW) Huang

Department of Electrical and Computer Engineering

University of Utah, Salt Lake City, UT



CPU Parallelism using C++ Thread

- Creating a `std::thread`
- Caring about const- and reference-semantics (`std::cref` and `std::ref`)
- Solving race conditions (`std::mutex`)
- Getting return values (`std::future<T>`)
- Allowing threads to complete / cleanup (`std::thread::join` or `std::thread::detach`)
- General Algorithm: partitioning data
- Style/Cleanup – Using C++ Lambdas

Creating a `std::thread`

```
#include <thread>
using std::thread;
```

```
thread (fun, args...)
```

- **fun** - a function we wish to have a thread run
- **args...** - any number of arguments we wish to pass to fun

<https://en.cppreference.com/w/cpp/thread/thread/thread>

```
thread t1(sprintf, "Hello from other thread", 1);
thread t2(vector<int>::push_back, v, 1);           // calls v.push_back(1)
```

Caring about const- and Reference-Semantics

- Functions can have varying signature types
- It can be very difficult to disambiguate between function overloads
- Threads can **only** be constructed with value-semantics
- Idea: Introduce a reference wrapper type which encapsulates references and const-references
 - <https://en.cppreference.com/w/cpp/utility/functional/ref>
 - `#include <functional>`
 - `std::cref` for const-reference
 - `std::ref` for reference

std::cref and std::ref

```
#include <functional>
using std::ref;
using std::cref;
void myFunction(const vector<int>& v, int& result);
...
// Usage:
thread t(myFunction, cref(vec), ref(ans));
```

Not Using Reference Wrappers...

```
/usr/include/c++/9.2.0/thread: In instantiation of 'std::thread::thread(_Callable&&, _Args&& ...) [with _Callable = void (&)(int&); _Args = {int&}; <template-parameter-1-3> = void]':
example.cpp:13:20:   required from here
/usr/include/c++/9.2.0/thread:120:44: error: static assertion failed: std::thread arguments must be invocable after conversion to rvalues
   120 |         typename decay<_Args>::type...>::value,
       |                                     ^~~~~
/usr/include/c++/9.2.0/thread: In instantiation of 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >':
/usr/include/c++/9.2.0/thread:131:22:   required from 'std::thread::thread(_Callable&&, _Args&& ...) [with _Callable = void (&)(int&); _Args = {int&}; <template-parameter-1-3> = void]':
example.cpp:13:20:   required from here
/usr/include/c++/9.2.0/thread:243:4: error: no type named 'type' in 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >::__result<std::tuple<void (*)(int&), int> >':
   243 |     _M_invoke(_Index_tuple<Ind...>)
       |     ^~~~~~
/usr/include/c++/9.2.0/thread:247:2: error: no type named 'type' in 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >::__result<std::tuple<void (*)(int&), int> >':
   247 |     operator()()
       |     ^~~~~~
make: *** [builtin]: example] Error 1
```

Not Using Reference Wrappers...

```
/usr/include/c++/9.2.0/thread: In instantiation of 'std::thread::thread(_Callable&&, _Args&& ...) [with _Callable = void (&)(int&); _Args = {int&}; <template-parameter-1-3> = void]':
example.cpp:13:20:   required from here
/usr/include/c++/9.2.0/thread:120:44: error: static assertion failed: std::thread arguments must be invocable after conversion to rvalues
   120 |         typename decay<_Args>::type...>::value,
       |                                     ^~~~~
/usr/include/c++/9.2.0/thread: In instantiation of 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >':
/usr/include/c++/9.2.0/thread:131:22:   required from 'std::thread::thread(_Callable&&, _Args&& ...) [with _Callable = void (&)(int&); _Args = {int&}; <template-parameter-1-3> = void]':
example.cpp:13:20:   required from here
/usr/include/c++/9.2.0/thread:243:4: error: no type named 'type' in 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >::__result<std::tuple<void (*)(int&), int> >':
   243 |     _M_invoke(_Index_tuple<Ind...>)
       |     ^~~~~~
/usr/include/c++/9.2.0/thread:247:2: error: no type named 'type' in 'struct std::thread::_Invoker<std::tuple<void (*)(int&), int> >::__result<std::tuple<void (*)(int&), int> >':
   247 |     operator()()
       |     ^~~~~~
make: *** [builtin]: example] Error 1
```

Solving Race Conditions

- Consider the following:

```
thread t1(std::vector<int>::push_back, ref(v), 1);  
thread t2(std::vector<int>::push_back, ref(v), 2);  
v.push_back(3);  
t1.join();  
t2.join();
```

- "v" can be modified **concurrently** by t1, t2, or the "main" thread!
- push_back() is **NOT** thread-safe!

Solving Race Conditions

- Idea: only allow **one thread** in a region at a time
 - Term: mutual exclusion
- In programming, we use a special object to represent this idea – mutex
 - **Mutex** = **mutual exclusion**
- Mutexes can “lock” and “unlock”

```
#include <mutex>  
using std::mutex;
```

std::mutex

```
// m is shared among all threads
mutex m;
// This is a place where multiple threads can be

m.lock(); // begin mutex region
// only ONE thread is allowed in here at a time
m.unlock(); // end of mutex region
```

Making a `std::mutex` shared

- Option 1: pass as a parameter to a “thread-safe function”

```
template <typename T>
Void safe_push_back(vector<T>& v, const T& value, mutex& lock)
{
    lock.lock();
    v.push_back(value);
    lock.unlock();
}
```

- Option 2: make it global
 - LOL JK **DON'T EVER DO THIS**

Making a `std::mutex` shared

- Calling the safe function:

```
vector<int> v = ...;
mutex m;

thread t1(safe_push_back<int>, ref(v), 1, ref(m));
thread t2(safe_push_back<int>, ref(v), 2, ref(m));
safe_push_back(v, 3, m);
t1.join();
t2.join();
```

Getting Return Values

- Up until this point, threads could call functions but not “return” anything
- We have functions that should return a value!
- Idea: threads run and take some time – we will get its result in the **future**
- Solution: introduce the concept of a future

```
#include <future>  
using std::future;
```

Creating a `std::future`

- There are a few ways in C++ to create a future, but we will only focus on ONE
 - `std::thread` allowed us to create a thread which we know would run asynchronously to our main thread
 - We want to asynchronously run and get a **future**
 - Introducing: `std::async`

```
using std::async;
```

```
// similar to thread
```

```
async (std::launch::async, Fun, Args...)
```

std::future – “get()”-ing the result

```
int rand_between (int low, int high) {  
    static std::minstd_rand rng{0};  
    return low + rng () % (high - low);  
}
```

```
future<int> result = async (  
    std::launch::async, rand_between, 0, 10);  
cout << result.get() << endl;
```

Managing ALL THE THREADS

Usually we will have some procedure that does the following:

1. Create a bunch of threads
2. Assign them to do a piece of the work
3. Wait for them to finish
4. ???
5. Profit

Managing ALL THE THREADS – CREATION

- Creating a bunch of threads
 - And have them do a work

```
vector<thread> threads;  
  
for (int tid = 0; tid < numThreads; ++tid) {  
    threads.emplace_back ( /* args to thread ctor */ );  
}
```

Managing ALL THE FUTURES – Wait/Cleanup

- Wait for them to finish

```
for (auto& t : threads) {  
    t.join();  
}
```


Managing FUTURES – Wait/Cleanup

- Wait for them to finish

```
for (auto& t : threads) {  
    auto result = t.get();  
    // do something with result  
}
```

Partitioning Data

Given a range of values from [low, high)

- We would like to be able to divide the work
- Work should be block-distributed
- Work should be evenly-divided

- **Goal: Do this efficiently**
- **Strategy: Consider “size”, number of “workers”, and current “worker ID”**

Partitioning Data

- A “range” is defined as a lower-bound and upper-bound
 - Initially, this is often $[low, high)$ or $[0, N)$
- When we have two workers, we want to partition as such:
 - Worker 0: $[0, N/2)$
 - Worker 1: $[N/2, N)$
- When we have three workers...
 - Worker 0: $[0, N/3)$
 - Worker 1: $[N/3, 2*N/3)$
 - Worker 2: $[2*N/3, N)$

Partitioning Data

```
int get_index (int elems, int id, int total)
{
    // ideally: elems * (id / total)
    return static_cast<long long>(elems) * id / total;
}
```

Partitioning Data

- Our ranges can be generalized... for N elements with P workers:

- **Worker 0:**

- Low: `get_index(N, 0, P)`
- High: `get_index(N, 1, P)`

- **Worker 1:**

- Low: `get_index(N, 1, P)`
- High: `get_index(N, 2, P)`

- **Worker $P-1$:**

- Low: `get_index(N, $P - 1$, P)`
- High: `get_index(N, P, P)`

Worker i :

- Low: `get_index(N, i , P)`
- High: `get_index(N, $i + 1$, P)`

Using C++ Lambdas

- Consider the following scenario:
 - You have a vector of ints you'd like to populate
 - You want to distribute the work across threads
 - You want each thread to add K copies of itself to the vector

Definition:

```
void populate (vector<int>& v, mutex& m, int K, int tid);  
// v, m, and K are the same for ALL  
// have to use ref() for v and m
```

Usage:

```
thread t1 (populate, ref(vec), ref(mut), K, tid);
```

Using C++ Lambdas

```
// vec, mut, and K are actual parameters
auto populate = [&vec, &mut, K] (int tid) {
    for (int i = 0; i < K; ++i) {
        mut.lock();
        vec.push_back(tid);
        mut.unlock();
    }
};
```

Using C++ Lambdas

```
vector<int> vec;  
mutex mut;  
int K = ...;  
auto populate = ... /* lambda definition */;  
for (int tid = 0; tid < P; ++tid) {  
    threads.emplace_back (populate, tid);  
}
```

Using C++ Lambdas

- Before

```
void populate (vector<int>& v, mutex& m, int K, int tid);  
threads.emplace_back(populate, ref(vec), ref(mut), K, tid);
```

- After

```
auto pop = [&vec, &mut, K] (int tid) {  
    return populate (vec, mut, K, tid);  
};  
threads.emplace_back(pop, tid);
```